

Roy (Yuanxi) Cheng

Laguna Niguel, CA | (949) 378-0596 | roycheng27@ucla.edu | [linkedin.com/in/ycheng05](https://www.linkedin.com/in/ycheng05)

EDUCATION

University of California, Los Angeles (UCLA)

Sep 2023 – Mar 2027

Bachelor of Science in Statistics and Data Science, Minor in Data Science Engineering

GPA: 3.75

- **Awards & Programs:** Dean's Honor's List, Persolv AI Intensive, BCG Case Comp Semi-finalist
- **Certificates:** Data Science & SQL in Python, EDA in R, Data Manipulation with dplyr, Data Visualization with ggplot2
- **Relevant Coursework:** Data Structures & Algorithms, Machine Learning, Neural Networks & Deep Learning, Data Analysis & Regression, Data Mining, Probability & Statistical Inference, Databases, NLP, Monte Carlo, Optimization

PROJECTS AND RESEARCH

Biomedical Tissue Imaging | *Python, PyTorch, YOLOv8, Pandas, NumPy, Matplotlib, OpenCV*

- Published article evaluating CNN-based instance segmentation architectures with AP metrics, benchmarking cytoplasm boundary precision for biomedical imaging
- Converted 4k+ CytoNuke annotations to COCO format through a custom data pipeline, enabling standardized training across 3 segmentation architectures
- Trained and tuned YOLOv8 instance segmentation models in PyTorch, improving cytoplasm AP50 by 5% over the Cyto-RCNN baseline

COVID Pandemic Effects on Secondary Education | *Python, Pandas, Tableau, Excel, Matplotlib*

- Cleaned and merged 100k+ student records across multi-level datasets with pandas, enabling accurate longitudinal analysis in Tableau
- Aggregated 4 years of test scores into time series visualizations, revealing a 15% average decline in student performance
- Created an interactive Tableau map linking geographic data to score changes across 50 U.S. states for a public research website

Adrenaline and Cognitive Fatigue Effects on Athletic Performance | *R, ggplot2, tidyverse, dplyr*

- Designed 2^3 factorial experiment with 96 participants across 8 conditions, blocking by age and gender to reduce variance
- Conducted three-way ANOVA and Tukey post hoc tests in R, revealing no significant interactions at $\alpha = 0.05$
- Validated ANOVA assumptions using log transformation and ggplot2 diagnostics, confirming normality and homoscedasticity of residuals

WORK AND LEADERSHIP EXPERIENCE

Capital One

McLean, Virginia

Product Management and Analysis Intern

Jun 2026 – Present

- Analyze corporate-wide people data with SQL and Python, identifying factors tied to successful organizational transitions across business units
- Synthesize findings into a mock HRBP dashboard of key reorg metrics, presenting to SVP to help HRBPs track retention, mobility, and headcount during reorganizations
- Architect a 6-stage Python pipeline ingesting 1,459 Snowflake performance reviews through PII anonymization, BGE semantic filtering, and FAISS retrieval before routing to an LLM DWA scoring agent
- Build a custom LLM client with sliding-window token rate limiting and retry-aware backoff to maintain reliable multi-hour inference throughput across the production review-processing pipeline

Data Science Union

Los Angeles, California

External Vice President

Dec 2025 – Present

- Directed external strategy across client partnerships, marketing, and professional development, expanding consulting and career opportunities for members
- Managed 3 directors and coordinated executive board priorities, aligning cross-functional execution to deliver quarterly and annual objectives
- Cultivated industry and alumni partnerships, securing client data science projects and expanding career programming across the organization

Kohl's

Los Angeles, California

Project Manager and Data Analyst

Jun 2025 – Mar 2026

- Developed and deployed a two-stage XGBoost forecasting system with time-aware validation to prevent leakage, cutting shipment arrival MAE by 83%
- Engineered route, carrier, seasonality, and port congestion features in Vertex AI Workbench, capturing delay patterns that improved transit time predictions

- Aggregated 8.5M+ product-level logistics records into shipment-level training data with SQL and Python, eliminating duplication and leakage-prone fields

ACTUM Digital

Prague, Czech Republic

Data Analyst and AI Intern

Jun 2025 – Aug 2025

- Built agentic Python pipelines to filter documents, generate metadata, and embed content into a Milvus vector database for retrieval-augmented workflows
- Created internal AI assistant by designing prompts and integrating an MCP server, co-presenting the live demo to ACTUM executives
- Automated extraction of structured data from 600+ unstructured PDFs with Python, consolidating results into a centralized dataset for client reporting

Million Dollar Baby

Los Angeles, California

Data Analysis Consultant

Apr 2025 – Jun 2025

- Analyzed \$10M+ in advertising spend with SQL, surfacing efficiency insights that informed growth strategy and marketing investment decisions
- Presented CAC and RoAS assessments across multiple attribution windows, guiding stakeholder decisions on marketing spend allocation
- Built an Excel dashboard tracking 20+ marketing KPIs, giving brand managers a recurring view of acquisition and return trends

TECHNICAL SKILLS

Languages: Python, C++, Java, SQL, R, JavaScript

Libraries/Frameworks: pandas, NumPy, PyTorch, scikit-learn, Matplotlib, TensorFlow, Keras, XGBoost, LangChain, MLflow

Tools & Platforms: Git, GitHub, AWS, GCP, Claude Code, BigQuery, Snowflake, Vertex AI, Tableau, Jupyter Notebooks